

Implementing a Secure Microservice Architecture: NGINX & Kubernetes Deployment

Orchestration: Alexander Moomaw
Eastern Washington University

Logging: Brendan Hopkins
Eastern Washington University

Defense: Cameron Olivier
Eastern Washington University

Services & Background: Andree Ramirez
Eastern Washington University

Frontend: Beighlor Martinez
Eastern Washington University

Services & Background: Dillon Pikulik
Eastern Washington University

Project Management: Chelsea Edwards
Eastern Washington University

Abstract

This paper presents a secure and scalable deployment architecture for web server environments using Kubernetes. The deployment incorporates a proxy and web application firewall for enhanced security and modular design. Additionally, we introduce our solution to logging and IP banning mechanisms to monitor and restrict unauthorized access. Our approach ensures high availability, load balancing, and robust defense against threats, offering a comprehensive framework for modern web applications.

1 Introduction

Ensuring the security and scalability of web server environments is a challenge in modern digital infrastructure. As cyber threats become more advanced, it's crucial to implement solutions that protect against unauthorized access while also maintaining high availability and performance.

To address these challenges, we propose a comprehensive deployment architecture by first utilizing Kubernetes, an open-source platform designed for automating deployment, scaling, and management of containerized applications [9]. With this platform, we ensure high-availability to prospective users and a modular environment to build and scale upon as we please.

To enhance our microservice architecture, we introduced an additional abstraction layer using an NGINX proxy hosted locally on the server. The proxy server fetches content from our Kubernetes load balancer, serving as a gateway between our internal services and external communications. To mitigate potential threats, we integrated ModSecurity, a well-known web application firewall (WAF), into our proxy [11]. ModSecurity monitors and blocks common exploitation methods by referencing a comprehensive rule list that addresses specific exploits included in the OWASP Top Ten [7]. Exploits that are specifically blocked include SQL injections, cross-site scripting, and command injection [8].

Exploiting our endpoint services isn't the only attack vector within our environment. To address this, we implemented

host-level logging to monitor network traffic and identify IP addresses not communicating via SSH or HTTP. Additionally, we applied an automated process to block malicious IP addresses, enhancing our overall security posture.

2 Background

Going into this project we had some inspiration. Back in the class CSCD 330 in Eastern Washington University (EWU), taught by Tony Antonio Espinoza, we had an assignment to get the weather from the geo location of a website. We use different tools and languages to do this. This mainly inspired the creation of the services part of this project.

The tool that we used from the inception is Flask. We used Flask to host our website. Flask is a lightweight and flexible Python web framework that allows for quick development of web applications. It provides essential tools because it runs in Python. All the services and the entirety of the website is hosted on Flask.

The Flask website is hosted in a Docker container. Docker is a platform that enables developers to automate the deployment of applications inside lightweight, portable containers. We used Docker to host the Flask website, ensuring a consistent environment and simplifying the deployment process.

We use `tcpdump` which is a powerful command-line packet analyzer used for network troubleshooting and analysis. It captures and displays network traffic, providing detailed insights into data packets transmitted over a network in real-time. We use `tcpdump` to watch what is happening on our network and looking for any suspicious activity which we will blocked from accessing the website.

The way we block user is with `iptables`. In `iptables` you can set rules. We use this to set rules to block the IP address that are involved with the suspicious activity that we detected using `tcpdump`.

We also used tools called `cronjobs` and `crontab`. `cronjobs` are scheduled tasks in Unix-like operating systems that run at specified times or intervals, automating repetitive processes. These tasks are managed using the `crontab` file,

where users can define the timing and commands for each job. We used these tools to do many task, which needed automation.

To address the challenges of securing and scaling web server environments, we utilized Kubernetes, an open-source platform for automating the deployment, scaling, and management of containerized applications. Kubernetes is widely adopted for its ability to provide high availability and a modular environment conducive to efficient application building and scaling.

To further enhance our microservice architecture, we integrated NGINX as a reverse proxy. NGINX directs traffic between clients and servers, adding an extra layer of security by handling requests before they reach the internal services. Additionally, we implemented ModSecurity, a web application firewall integrated with NGINX. ModSecurity is effective in detecting and mitigating common web application attacks, thereby strengthening our security measures.

3 Orchestration

The ultimate challenge of any deployment life cycle is tying all ideas into a single, nicely wrapped package. This is the goal of orchestration. To implement, or *orchestrate* a deployment such as this, we designed a multi-tiered architectural diagram (Figure 1) that outlines how every building block in our environment fits together to form a finalized representation of a working project.

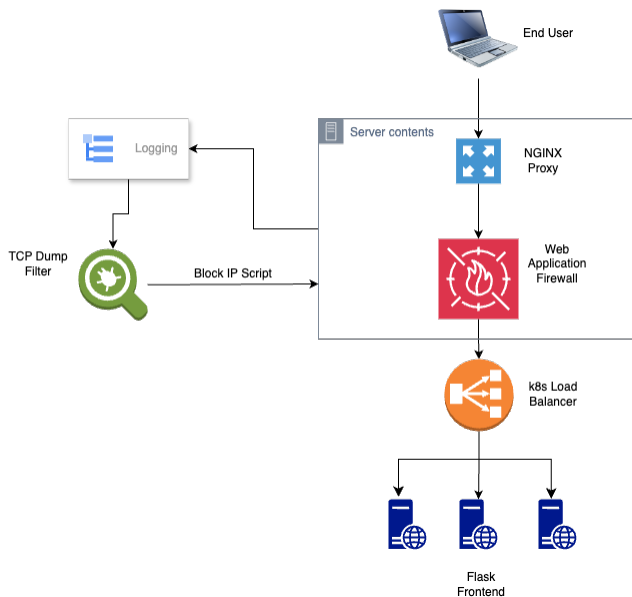


Figure 1: Multi-tiered architectural diagram.

To host our content, we utilized Kubernetes. This method first involved integrating a docker image that was specifically designed to host content created by our frontend develop-

ers. We then created service and deployment YAML files to control multiple containers behind a load balancer.

Once we had a functional cluster of orchestrated containers, the next goal was to make our load balancer not the first point of contact between external users and our internal environment. This involved integrating an NGINX reverse proxy [6] on our host machine to facilitate communication without exposing our services directly to the internet.

To setup an NGINX reverse proxy, we had to install the service then modify the global configuration file in the `/etc/nginx/conf.d/` directory to specifically mediate communications between external users and our internal load balancer [12]. The specifics for this step is very convoluted, essentially we had to define server and location blocks with a `proxy pass` flag to enable the proxy feature of NGINX.

To tie our scalable, secure environment together on the orchestration end, we installed a web application firewall (WAF) on our reverse proxy to inspect requests being sent to the load balancer [10]. The requests will be meticulously cross-referenced against a series of rule lists specifically designed to catch exploits matching the OWASP Top Ten vulnerabilities. If a exploit request is detected, the WAF will drop the request and return a HTTP 403 status code to the threat actor.

4 Logging and IP Gathering

To enhance the security of our web server environment, we implemented a robust mechanism for gathering and logging IP addresses. This approach involves capturing IP addresses attempting to access our server on ports other than 22 (SSH) and 80 (HTTP) and maintaining comprehensive logs for further analysis. The IP gathering mechanism is designed to identify unauthorized access attempts by capturing IP addresses that try to connect to non-standard ports. This is achieved through a script that utilizes tcpdump to monitor network traffic and filter out unwanted connections [5]. The command

```
tcpdump -i any "ip and not host $my_ip and not po
```

will listen for all connections not on ports 22 and 80, and export them to a temporary file. This is because those are the ports my team will use to connect and test the website, and we don't want to get blocked. The command

```
(grep -oP '(\d{1,3}\.){3}\d{1,3}' "$temp_file" |
```

searches the temporary file for any IP addresses and exports them to the IP list. The captured IP addresses are stored in a cumulative list, which can be used to block malicious IPs via iptables. These commands are executed every two minutes via a cron job, which reduces the attackers' window of opportunity and ensures that the temporary file remains manageable in size.

In addition to gathering IP addresses, we implemented an hourly logging mechanism to maintain detailed records of network activity. By excluding traffic on ports 22 and 80, we

focus our monitoring on potentially unwanted or malicious attempts. This involves running a script that logs all other IP traffic, storing the data with timestamps [13]. These logs are crucial for analyzing traffic patterns and identifying potential security threats. The logging process is automated with a cron job that runs every hour. This schedule helps us pinpoint the peak times attackers are trying to breach our defenses and allows us to trace the method and timing of any successful intrusions. The logging gives us the port number they accessed and the IP of the attacker. With this information, we can filter that port and attempt to block them from getting in that way again. To ensure no data is lost during log creation, the `tcpdump` process is stopped two minutes before the end of each hour, giving the script enough time to write to the timestamped log file. To keep our logging effective, we regularly review the logs. These reviews involve checking the log files to ensure they are accurately capturing data as intended. We also analyze the recorded attempts to understand the strategies attackers are using to target our server. This ongoing evaluation helps us stay informed about new threats, enabling us to adapt our defenses accordingly.

5 Defense

We have established a comprehensive security protocol to monitor and block unauthorized access attempts to our web server. This involves capturing IP addresses attempting to access our server on ports other than 22 (SSH) and 80 (HTTP) and maintaining comprehensive logs for further analysis.

The process begins by clearing the log file `/var/log/honeypot/block_ip.log` to ensure new entries are recorded clearly.

The script then checks for the existence of a cumulative IP list file at `/var/log/honeypot/ip_list_all.txt`, which contains IP addresses flagged for suspicious activity. If this list exists, the script reads each IP address and adds a `DROP` rule to `iptables`, effectively blocking incoming traffic from these addresses. The success or failure of each attempt to block an IP address is logged with a timestamp in the `block_ip.log` file.

If the cumulative IP list is not found, an error message is logged and the script exits.

6 Threat Model

The ‘Weather and Jokes’ web application employs Kubernetes and NGINX within its architecture to provide defenses against several cybersecurity threats. We are specifically defending against unauthorized access, SQL injections, Cross-Site Scripting (XSS), session hijacking, and Denial of Service (DoS) attacks. In contemplating the perceived capabilities of potential attackers, we recognize that they are likely to possess a high degree of technical sophistication, equipped

with the skills necessary to exploit vulnerabilities typical of web applications. These attackers might use reconnaissance tools to gain information, and advanced techniques for SQL injections, craft XSS attacks, and leverage automated tools to execute DoS attacks.

To counter these threats, our application utilizes a load balancer within the Kubernetes environment to effectively manage traffic volumes that could potentially lead to DoS attacks. This load balancer helps distribute incoming traffic evenly across available servers, preventing any single server from becoming overwhelmed. We sanitize user input, protecting against XSS, and logging the attempts they made so we understand more of our attackers’ attacks and motives. Additionally, NGINX, configured as a reverse proxy, is important in safeguarding against unauthorized access and filtering out malicious requests. It acts as a gatekeeper, routing all incoming traffic through its server and using ModSecurity—a web application firewall (talked more about later)—to identify and block threats identified from patterns typical in SQL injections and XSS. Moreover, NGINX simulates secure user sessions by managing encrypted connections, a key tactic to make the honeypot appear as a legitimate and secure environment, enticing attackers to engage more deeply. These encrypted channels are designed to ensure that any data captured, including session tokens, remains controlled and uncompromised, allowing us to safely monitor and analyze malicious activities without risk. Through the strategic implementation of Kubernetes and NGINX, our application is well-prepared to defend against a wide spectrum of cybersecurity threats.

7 Frontend Design

We implemented our website using a Dockerfile pulling an image of Ubuntu which installs Flask that we use as our webserver. You can find more information on backend code and the Flask Server in the services section. We decided to use Flask based on previous experience with this service and Docker so we could containerize our project for added defense. By adding Docker we provided isolation from a singular host and allowed us to add defensive features without effecting our entire system. Our Website is called ‘Weather and Jokes’, and includes four pages: Home, Contact, FTP Manual, and Login. The Home page runs two API services, ‘Get a Joke’ and ‘Weather Checking’, the Contact page lists each of our team members and their roles within the group, the FTP Manual page gives instructions to our team members on how to use FTP and the final page, Login, takes in user input and redirects back to the login page.

The FTP Manual page is meant to be a distraction for threat actors and is not fundamental to the actual Website, but rather a ploy to try and catch attackers attempting to use FTP. The login page never actually connects to any backend service and is also used to try and catch potential threat actors attempting to use attacks such as SQL Injections. These features allow

our webpage to act more as a honeypot gathering information on potential attackers, and allowing our defense team to monitor/log their activities and if necessary block their IP address.

8 Services

We have 2 services, both run in `flask_example.py` file and are presented on the `index.html` page. Flask is the resource we used to host the website locally. Flask is part of Python, which is easy to work with and implement, that was the main reason for choosing Flask. One of our services is utilizing a weather API. This API has the weather in different geo locations around the world. We use this API to find the weather of a website's geo location. The user has a `textbox` on the `index.html` page that they can enter the website's URL. After they enter the URL they can hit the `button` to get the weather. We ask for user input of a website's URL and output the current weather of that website. This is done by taking the URL and running `gethostbyname` to get the IP address of that URL. Now after we get the IP address, we run a `whois` command on the IP address. This command will give us a lot of information on the IP address, but we parse out the physical address associated with the IP. After the physical address is parsed and formatted for the API we call (`https://geocoding.geo.census.gov/`) API and get returned the `X,Y` geo coordinates of that address [2]. Now that we have geo coordinates we can send those to the weather API. The weather API is at `https://api.weather.gov/points/`, where we input the coordinates at the end of the URL to get back the weather [3]. After we send the coordinates, we receive the weather of that location, which is the location of the initially given website. All of these commands are automated in the `flask_example.py` file. I chose to implement this service because I have done it in the past and had very well-documented files that aided in quick implementation. This in return aided the rest of the team by speeding up the process of deployment. This service also is very large and can be used for about every website, this could keep the attacks busy, possibly giving information away about themselves.

Another API service that was utilized was the `/joke` route in the Flask application serving a random dad joke to the user. When accessed, the function `joke()` is executed, which sets the API URL `https://icanhazdadjoke.com/` and defines a `cURL` command to request a joke in plain text format [1]. The `subprocess.run` function runs this command and captures the output, which contains the joke. Finally, `result.stdout` is returned, sending the joke back to the user's browser. Additionally, there is a clickable button in `index.html` that allows users to retrieve a joke, providing a simple and interactive way to display a random dad joke. The reason this service was chosen was to have clickable and interactable features on the websites so that targets on

the website are encouraged to click more and give potential information.

9 Evaluation Methodology

To adequately test our infrastructure, we set out to exploit the specific methods that the web application firewall was built to defend against. This is done through the login page of our frontend environment. A successful evaluation will be represented by any attack execution being redirected via a HTTP 403 response, indicating that the attack was successfully detected and blocked by our web application firewall.

The first exploit we chose was an SQL injection, specifically a tautology attack. This attack aims to break the conditional statement for a typical user name/password login combination on an SQL backend server such that when the query is made through the login page to check if the user exists the condition will always be true, thereby allowing the attacker to log in as any user that exists [14].

The second exploit we chose was a cross-site scripting (XSS) attack. This attack attempts to inject malicious scripts into input sections or source code of a website [15]. This attack will also be tested by appending an exploit to the end of our URL. If the threat actor is able to successfully execute their attack, they would be able to run arbitrary scripts in the context of the user's browser. This could lead to a range of malicious activities, such as stealing cookies, session tokens, or other sensitive information, defacing the website, or redirecting users to malicious sites. Our testing will involve various payloads to simulate real-world XSS attack vectors and assess the vulnerability of our web application to these types of exploits.

The final exploit we will be testing is command injection. This exploit attempts to bypass a input field in such a way that allows the attacker to execute arbitrary operating system commands on the server [16]. If an attacker is able to successfully perform this attack, they can compromise the web server, exfiltrate data, and establish persistence by injecting a backdoor into our system [4].

10 Results

The first test run was an SQL Injection on the Login page, using the expression " `admin' OR '1'='1' -` " in the user input with " `password` " in the password input. Normal behavior would just redirect the page back to itself, however because it detected an SQL Injection the page was forwarded to a 403 Forbidden page, and the suspicious activity was logged, with the time of attempted access, the attempted username, the attempted password, and the IP address of the attacker. The next attack we used was an SQL injection in the password input where user input was " `user` ", and password input was " `wrongpassword' OR 'a'='a` " again we got the same results as

the first attempt, it forwarded the client to a 403 forbidden page and the suspicious activity was logged.

The second test we ran was cross site scripting or XSS, so in normal behavior the pages would redirect normally, so 10.102.68.91/about to 10.102.68.91/home, however in testing we added an XSS script at the end of the url to test the behavior of our webpages. The command we used was `http://10.102.67.91/about=</script><script>alert(0)</script>`, an XSS script that would throw an alert on the webpage. However our defense worked as intended redirecting to a 403 forbidden page and did not create the XSS alert as the attacker intended. The suspicious activity was logged and given to our security team so they could block the IP address of the potential attacker. This 403 error actually gets thrown any time there is a reference using `</script>`. The normal activity for the webpage is to redirect to a url not found, for example if you enter 10.102.68.91/about=user it will redirect to URL not found, so any XSS or added script will be blocked by our firewall. We made sure to test each of our pages, home, about, manual, and login, and each page passed as intended.

The third test we ran was for command injection we again used on our login page. In normal behavior a successful login attempt would redirect back to the login page. The attack we tested command injection is a type of attack that takes advantage of a program using a operating system call, for example if our program ran ping and displayed results of the ping, an attacker would be able to input `'8.8.8.8;ls'` to display information about whats inside their current directory. The `';ls'` is the attack and 8.8.8.8 would be the normal input. So in our test we used the username input `'user;ls'` and password as `'password'`, the firewall successfully blocked this attack and redirected the page to a 403 forbidden page and logged the suspicious activity, alerting our monitoring team and indicating another successful test.

11 Conclusion

Our web server security solution deploys and scales with Kubernetes, supported by an NGINX proxy and ModSecurity for increased protection. Host-level logging and automated IP blocking further strengthen security. Using a multi-tiered architecture and Kubernetes hosting, we integrate and orchestrate these components seamlessly. An NGINX reverse proxy ensures secure communication while a web application firewall examines and blocks known vulnerabilities.

To enhance security, we implemented a mechanism that logs IP addresses trying to access non-standard ports, with hourly logs for network activity monitoring. We utilize Docker to run our website 'Weather and Jokes' with Ubuntu and Flask, providing isolated defensive features. The website includes four pages: Home, Contact, FTP Manual (honeypot), and Login (trap for attackers).

Two services in the `flask_example.py` file fetch weather

and dad jokes using APIs. The weather service converts URLs to IP addresses, retrieves physical addresses, and obtains weather based on geographical coordinates. The joke service offers an interactive feature for users.

In order to evaluate our infrastructure, we conduct simulated attacks, including SQL injection, cross-site scripting (XSS), and command injection, which our web application firewall successfully detects and blocks, as confirmed by HTTP 403 responses.

References

- [1] Dad jokes api. Available at: <https://icanhazdadjoke.com/>.
- [2] Geo location api. Available at: <https://geocoding.geo.census.gov/>.
- [3] Weather api. Available at: <https://api.weather.gov/points/>.
- [4] DRD. Use command injection to pop a reverse shell on a web server, 2018. <https://null-byte.wonderhowto.com/how-to/use-command-injection-pop-reverse-shell-web-server-01/>
- [5] Damon Garn. How to capture and analyze traffic with tcpdump, 2023. <https://www.techtarget.com/searchnetworking/tutorial/How-to-capture-and-analyze-traffic-with-tcpdump>.
- [6] Hostinger. nginx is an http and reverse proxy server. <https://www.hostinger.com/tutorials/how-to-set-up-nginx-reverse-proxy/>, 2024. Accessed: 2024-05-21.
- [7] <https://owasp.org/>. The owasp top 10 is a standard awareness document for developers and web application security. <https://owasp.org/www-project-top-ten/>, 2021. Accessed: 2024-05-21.
- [8] <https://owasp.org/>. Owasp modsecurity core rule set. <https://owasp.org/www-project-modsecurity-core-rule-set/>, 2024. Accessed: 2024-05-02.
- [9] Kubernetes. This page is an overview of kubernetes. <https://kubernetes.io/docs/concepts/overview/>, 2023. Accessed: 2024-05-21.
- [10] Linode. Securing nginx with modsecurity. <https://www.linode.com/docs/guides/securing-nginx-with-modsecurity/>, 2024. Accessed: 2024-05-02.
- [11] modsecurity. This page is an overview of modsecurity. <https://github.com/owasp-modsecurity/ModSecurity>, 2021. Accessed: 2024-05-21.

- [12] Nginx. How to set up nginx reverse proxy. <https://nginx.org/en/>, 2024. Accessed: 2024-05-21.
- [13] Daniel Schwartz. Schedule tcpdump with cron, 2018. Available at: <https://schwartzdaniel.com/schedule-tcpdump-with-cron/>.
- [14] Chandershekhar Sharma. Types of sql manipulation, 2016. https://www.researchgate.net/figure/Types-of-SQL-manipulation-a-Tautology-In-this-type-of-SQLIA-an-attacker-exploits-an_fig4_299575353.
- [15] Synopsys. Cross-site scripting (xss), 2022. [https://www.synopsys.com/glossary/what-is-cross-site-scripting.html#:~:text=Cross%2Dsite%20scripting%20\(XSS\)%20is%20an%20attack%20in%20which,the%20user%20to%20click%20it](https://www.synopsys.com/glossary/what-is-cross-site-scripting.html#:~:text=Cross%2Dsite%20scripting%20(XSS)%20is%20an%20attack%20in%20which,the%20user%20to%20click%20it).
- [16] Weilin Zhong. Command injection, 2022. https://owasp.org/www-community/attacks/Command_Injection#:~:text=Command%20injection%20is%20an%20attack.